

# Model Improvements

**Practical Machine Learning (with R)**

UC Berkeley

Spring 2016

# Topics

- Administrative
  - Role Call
  - Assignments due to github
- Review/Expectations
  - Readings
  - Previous Lecture
- New Topics



# REVIEW AND EXPECTATIONS





# RESAMPLING METHODS

- Get more accurate estimation of a statistic/value by resampling methods
- ➔ Improves function estimate



# Model Performance Evaluation

- Understand model performance metrics
- How they are calculated ...

## Classification

- Metrics: Accuracy, Kappa
- Graphical
  - `caret::confusionMatrix`
  - `pROC::ROC`

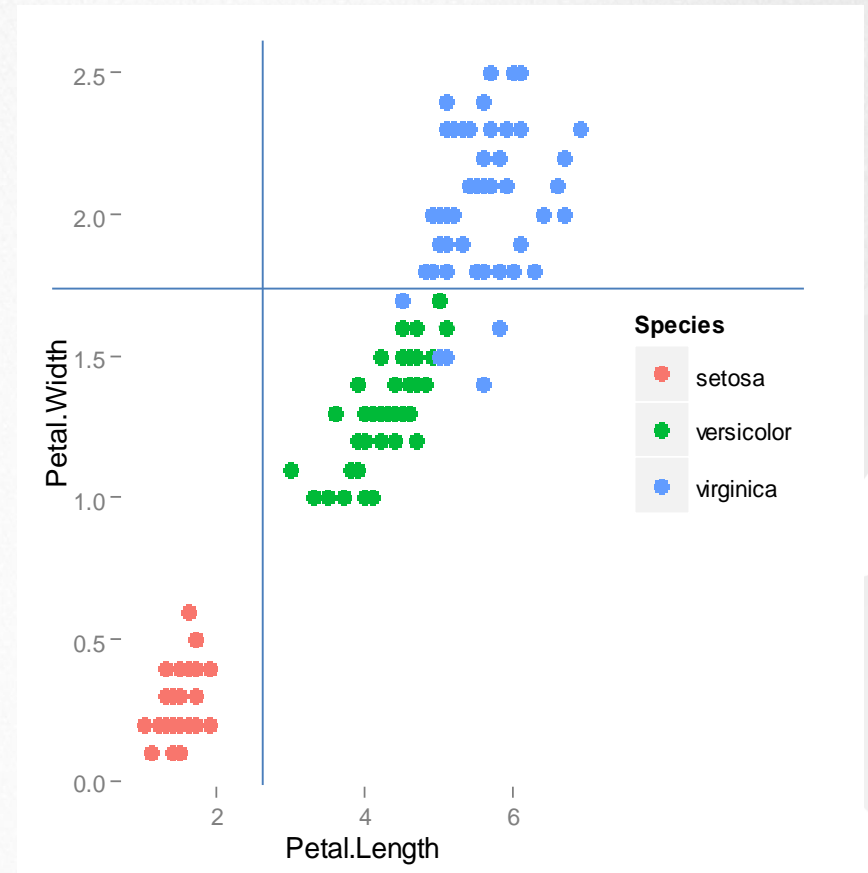
## Regression

- RMSE, MAE, etc.
- Graphical
  - Histogram of errors
  - $y$  vs  $\hat{y}$  scatterplot



# Decision Trees

- All about splitting
  - Different variants have different rules for evaluating the splits
- Tree = Ruleset = Partition of Space
  - Node = Rule = “box” (contiguous region of space)





# TREE VARIANTS

- ⇒ There are many tree variants
- ⇒ Tweaks
  - change how splits are determined
  - when to stop growing the tree
  - how the node value is determined



# RULES

- ➔ As derived from trees often have repeated conditions

```
NumCarbon > 3.777 &  
SurfaceArea1 > 0.978 &  
SurfaceArea1 > 8.404 &  
FP009 <= 0.5 &  
FP075 <= 0.5 &  
NumRotBonds > 1.498 &  
NumRotBonds > 1.701
```

M5rules, C4.5Rules

Rules and their conditions live on their own, conditions can be adjusted to help bias-variance trade-off



# MISSING DATA

- ➔ Missing values in predictors are common
- ➔ A split determines which observations go to the LHS and RHS. How to Handle `NA`s?
- ➔ `NA_Categorical`
  - Treat as separate category
- ➔ `NA` (in general)
  - Use **Surrogate Splits**



# SURROGATE SPLITS

- ⇒ Tree is built ignoring missing data
  - Any record with incomplete data (response or predictor) is rejected -or-
  - Missing data is not used for split-point determinations
- ⇒ Variables are often *collinear* → splits are similar and send variables down the same path.
  - Choose a surrogate (variable) split that best approximates the chosen split (accuracy)
  - Very often this is also a good split

# REVIEW AND EXPECTATION

## Understand **Recursive Partitioning** :

- intrinsically how recursive partitioning models work; how splits are determined; what splitting accomplishes
- how model tuning parameters control for bias variance trade-offs
- what  $C_p$  is and how to use it to prune trees to the proper size





# Tree Method Advantages

→ List em



# REVIEW AND EXPECTATION

Use **rpart** and/or **ctree**

- Build Recursive Partitioning Models using `rpart`
- Prune trees to statistically relevant size
- Plot `rpart` trees using
  - `party::as.party`
  - `map.tree::drawtree`



# UNDERSTAND CARET

- ➔ Use the **caret** package and the `train` function to build models
- ➔ Understand the difference between using `caret` and building models manually. What `caret` provides.
- ➔ Control how models are built using the `train` & `trainControl` functions
- ➔ How to extract the final model
- ➔ How to plot the tuning parameters, performance curves







# QUESTIONS



# IMPROVING MODELS



# TREE DISADVANTAGES

⇒ List em





# Tree Disadvantages

- Model instability (sensitive to data)
  - Derives from each subsequent split is dependent on prior splits
- Less than optimal predictive performance
  - Rectangular regions
- Limited number of outcome values  $\leq$  number of terminal nodes
- Intra-node high and low values not fit well
- Selection bias toward predictors with higher number of distinct values
- Tuning parameter,  $C_p$
- Splits of correlated variables ambiguous
- Treatment of missing values



# TWO BIG IDEAS

## ➤ **Wisdom of the crowds**

It is better to make estimates from multiple models (**ensembles**) than individual models

- Better predictions
- Lower variance for the same model

## ➤ **Greed is bad. Patience is good.**

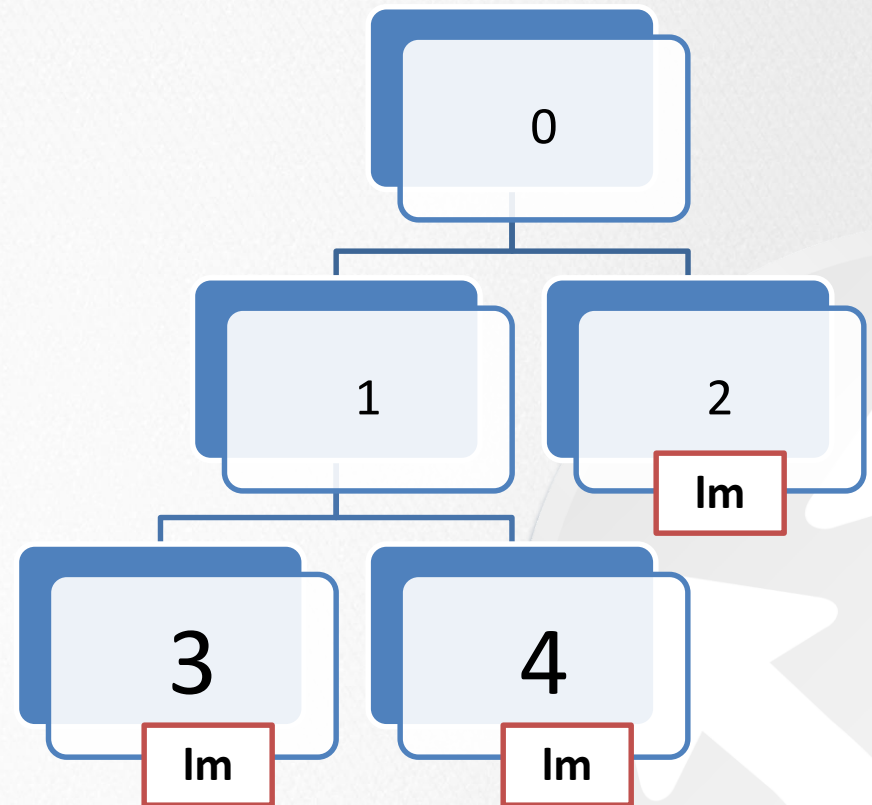
It is better to slowly approach your solution than arrive at an answer directly.

- More accurate solutions



# Tree Enhancement: M5

- **Wisdom of the Crowd!**
- Having one value represent the entirety of the node leaves information in the node.
- Function in the node is a simple average
- Use something better
  - **M5** put linear models in nodes of trees



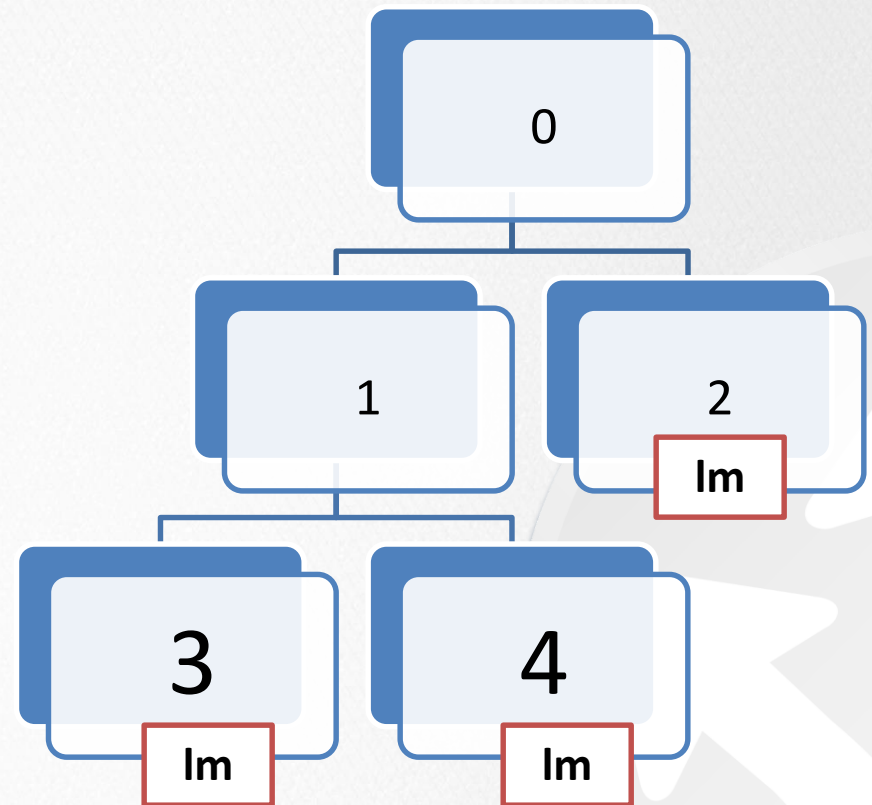
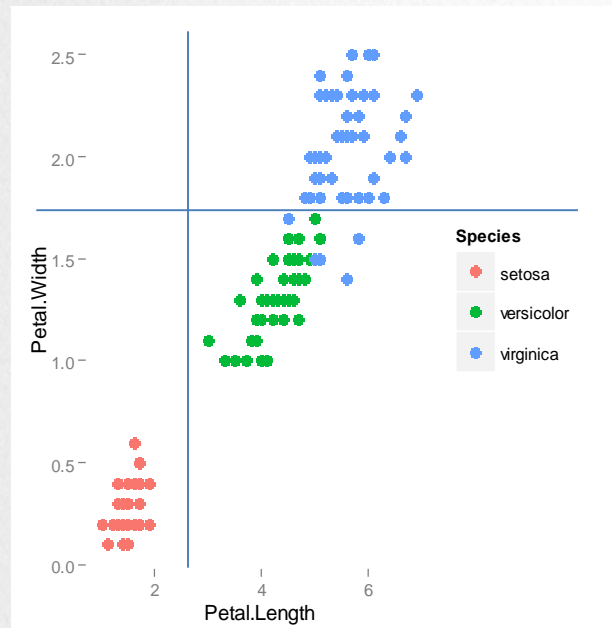


# M5 Tree Enhancement (cont.)

## → Greed is bad

- linear models are built on the residuals of the tree model.

- Models are recursive



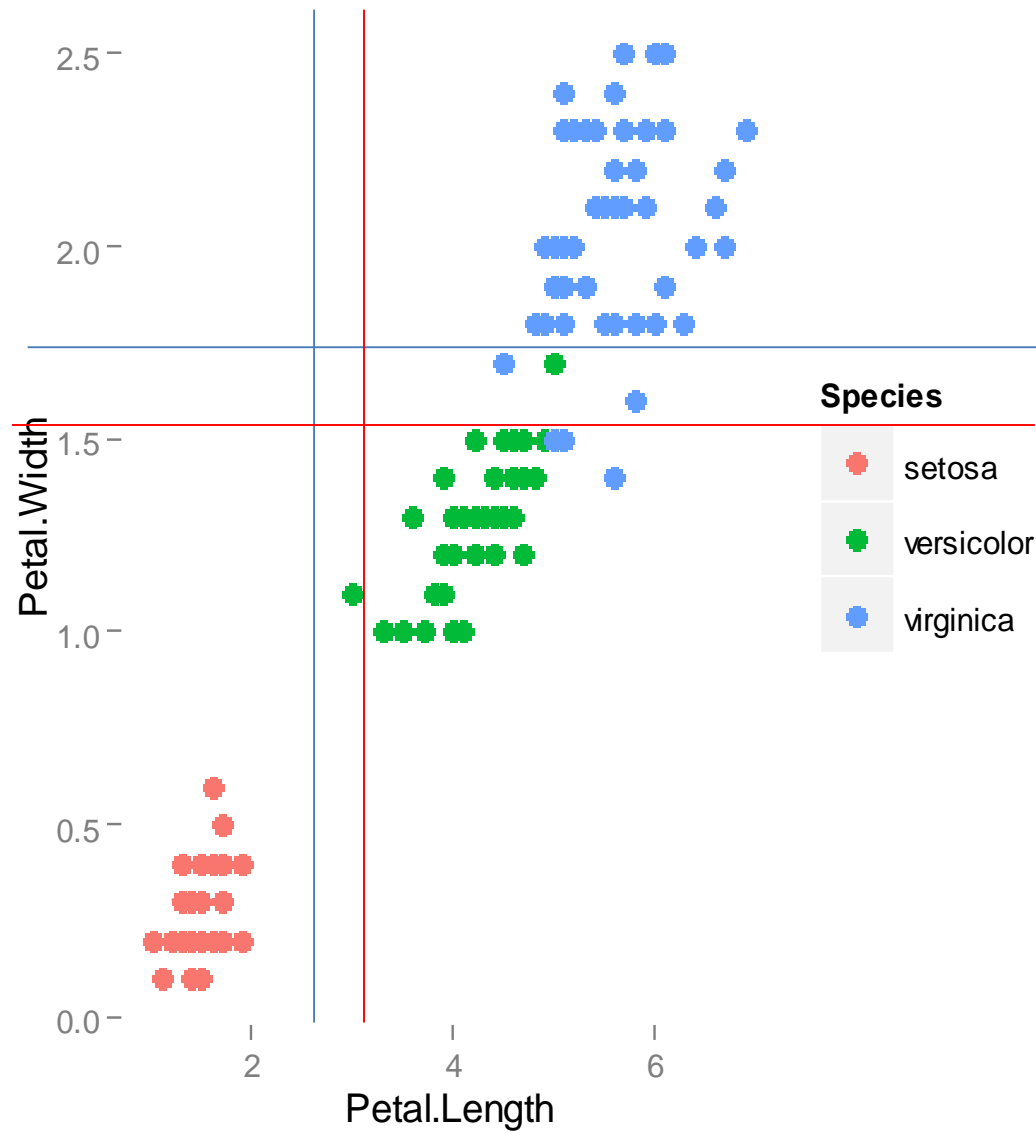
# BAGGING MODELS

➔ Brieman:

"Bagging is a general approach that uses bootstrapping in conjunction with any regression (or classification) model to construct an ensemble."

```
1 for  $i = 1$  to  $m$  do
2   |   Generate a bootstrap sample of the original data
3   |   Train an unpruned tree model on this sample
4 end
```

$$\hat{y} = \frac{\sum_i \hat{y}_i}{m}$$





# BAGGING NOTES

## ➤ Lowers variance

- Increases stability
- Has less effect on lower variance models (e.g. linear models)
- More effect on weak learners

## ➤ Disadvantages

- Computational cost → but parallelizable
- Reduces Interpretability



# RANDOM FOREST

- **Wisdom of the Crowds:** Bagging
- **Greed is bad:** consider subset of predictors at each split

```
1 Select the number of models to build,  $m$ 
2 for  $i = 1$  to  $m$  do
3     Generate a bootstrap sample of the original data
4     Train a tree model on this sample
5     for each split do
6         Randomly select  $k$  ( $< P$ ) of the original predictors
7         Select the best predictor among the  $k$  predictors and
           partition the data
8     end
9     Use typical tree model stopping criteria to determine when a
       tree is complete (but do not prune)
10 end
```

# TUNING PARAMETER

$m_{\text{try}}$  : number of predictors to use at each split

- **regression** 1/3rd of number predictors
  - **classification**  $\sqrt{\text{number of predictors}}$
- ➔ Kuhn: “Starting with five values of  $k$  that are somewhat evenly spaced across the range from 2 to  $P$ ”.





# ADVANTAGES

- ➔ No overfitting
- ➔ More trees better (limited by computation time/power only)
- ➔ In caret, parameters are considered independently
- ➔ Because each learner is selected independently of all previous learners, Random Forests is robust to a noisy response
- ➔ Computationally efficient -- each tree built on subset of predictors at each split.
- ➔ Use any tree variants as "base learner": CART, ctree, etc



# BOOSTING

⇒ Single models work;

- Multiple models work better

⇒ Idea is simple:

- **Fit first** model:  $\hat{y}_1 \sim f_1(x)$

- **Fit** errors/residuals: 
$$\begin{aligned}\hat{y}_2 &= f_2(y - \hat{y}_1) \\ &= f_2(y - f_1(x)) \\ &= f_2(x)\end{aligned}$$

- **Iterate:**  $\hat{y}_i = (y - \hat{y}_{i-1}) \sim f_i(x)$

- **Predict:**  $\hat{y} \sim \sum_i f_i(x)$



# BOOSTING NOTES

- ➔ Additive models
- ➔ Works best with “weak learners”
  - i.e. ungreedy, low bias, low variance
  - ~~Any~~ Most models with a tuning parameter can be a weak learner
  - Trees are excellent weak learners
    - Weak → “restricted depth”
- ➔ Residuals or errors define a gradient
- ➔ Interpreted as forward step-wise regression with exponential loss





# SIMPLE GRADIENT BOOSTING

- 1 Select tree depth,  $D$ , and number of iterations,  $K$
- 2 Compute the average response,  $\bar{y}$ , and use this as the initial predicted value for each sample
- 3 for  $k = 1$  to  $K$  do
  - 4 Compute the residual, the difference between the observed value and the *current* predicted value, for each sample
  - 5 Fit a regression tree of depth,  $D$ , using the residuals as the response
  - 6 Predict each sample using the regression tree fit in the previous step
  - 7 Update the predicted value of each sample by adding the previous iteration's predicted value to the predicted value generated in the previous step
- 8 end

# Simple Gradient Boosting – Comparison To Random Forest

Similarities

Differences



# STOCHASTIC GRADIENT BOOSTING

## ⇒ Gradient Boosting Susceptible to Overfitting

- Apply “regularization/shrinkage”
  - Use  $\lambda$  (“Learning Rate”)  
Rather than add the entirety of the residuals, add a fraction of the residuals at each iteration.

$$\hat{y} \sim \lambda \sum_i f_i(x) \quad 0 < \lambda \leq 1$$

- Small values for  $\lambda$  ( $\sim 0.01$ ) work best
- $\lambda \sim 1/\text{computational time} \sim 1/\text{storage time}$

## ⇒ Use bagging, as well

- Bagging Fraction: a sample of data in each loop iteration





# APPENDIX



